

Section 1: Who I am and what I came in wanting to do (30 seconds, context setting)

Spoken script / speaker notes:

My name is Polo. I came into this semester passionate about two things: filmmaking and storytelling. The craft of taking something raw and shaping it into something that actually lands with people. That is what drives me.

Before the semester started I went to San Francisco. I was staying around Mission and Dogpatch, right in the middle of where a lot of the startup activity happens. I went in knowing it would be interesting. I did not expect it to hit the way it did.

I was in rooms with founders who were unbelievably young, building things that most people in New Zealand have not even heard of yet. Every single startup had AI at the centre of it somehow, even the ones where it did not really fit. The word had become a badge that everything had to wear.

But the thing that really got me was a specific kind of person they call the doomers. Founders and operators who were actively shipping AI products while being quietly convinced it might not end well. Sitting in a room with someone building the very thing they were afraid of is a strange experience. But they were serious, and that seriousness stuck with me.

That tension is what generated everything I worked on this semester. Because the question underneath all of it is this: what does it actually look like when AI assists creative work rather than replaces it? And can you design something that makes people feel that difference rather than just understand it?

That is what I came in wanting to explore. Everything else this semester came from trying to answer that.

Key points to hit:

- Passionate about filmmaking and storytelling, that is the lens for everything
- San Francisco trip happened before semester, not during, it came in with me
- The doomers: founders building the thing they feared, that contradiction matters
- The central question: AI assisting creativity versus replacing it
- Everything this semester came from trying to answer that one question

Section 2: The local issue and brief I'm working within (addresses the DES300 question honestly)

A quick note before I get into the work. The rubric for this presentation asks about connections to a DES300 proposal. I completed DES300 last semester, so rather than forcing a connection that isn't there, I want to be honest about that and speak to what I can actually speak to: the local issue I identified through my own research, the brief I chose, and where this work points toward 304.

The brief I chose was Emerging Technologies, with the Create lens. The reason was straightforward. The Create lens is fundamentally about how emerging technology changes what it means to make things, be a creator, and own your work. That question sits inside everything I care about. It was the one brief that gave me genuine room to move in.

The local issue I kept coming back to was this: everyday creative practitioners in Aotearoa, filmmakers, musicians, designers, painters, are facing a version of the same question the doomers in San Francisco were already treating as settled. What happens to your creative identity when a tool can do more and more of what you used to do yourself? That conversation is already happening here. It is just quieter, and less visible.

What San Francisco gave me was a glimpse of where that conversation lands when it has had five more years to play out. And that became the raw material for everything I built this semester. A speculative future Aotearoa where that tension has already settled into something, and three different formats for putting an audience inside it.

The How Might We question I landed on was this: how might we use speculative design to help everyday creators in Aotearoa feel what it means to work alongside AI, rather than be replaced by it?

That question drove every experiment from Week 4 onwards.

Key points to hit:

- DES300 was completed last semester, being upfront about that rather than dodging it
- Chose Emerging Technologies stream, Create lens, because it gave the most room to move
- Local issue: everyday creatives in Aotearoa facing AI displacement, quieter version of what SF already treats as settled
- SF trip gave the raw material, speculative future Aotearoa is where it lands
- HMW question: how might we use speculative design to help everyday creators in Aotearoa feel what it means to work alongside AI, rather than be replaced by it
- Everything from Week 4 came from trying to answer that

Section 3: The research and references that shaped my thinking (SF trip, doomers, Dunne and Raby, Kirby)

Before I started making anything, I needed to understand what I was actually trying to make. Not just the format, but the feeling. The brief used a phrase that kept jumping out at me: placing an audience inside a speculative scenario. That is not just describing a concept. That is describing an experience. And it made me think about the work that had actually done that to me.

Two references shaped how I thought about format. The first was Dunne and Raby's *Speculative Everything*. Their argument is that speculative design is not about predicting the future. It is about using design to open up debate about how we want to live. The three things I built this semester are not answers. They are prompts. That framing came directly from them.

The second was David Kirby's concept of the diegetic prototype. A diegetic prototype is an artefact that appears to come from a possible future rather than describing that future from the outside. The distinction matters because it changes what the audience is being asked to do. A film about AI displacement asks you to watch and imagine. An artefact from inside that world drops you into it and asks you to respond. That became the test I held every experiment against.

On the media side, two things shaped how I thought about the interactive format specifically. *Emily is Away*, a game set entirely inside a fake AIM messenger interface, showed me that an interface itself can be a world. You are not watching someone have a conversation. You are inside one. And *Bandersnatch* showed me what it feels like to be an active participant rather than a passive viewer. Once you have felt that difference, linear content feels like it is missing something.

The *Cluely* launch video was a reference point for tone on the film side. It presents itself as if the product already exists and the question is just whether you are keeping up. That confidence of format is what I wanted the *Compass* film to sit in.

All of that research fed directly into the three experiments. It was not background reading. It was the scaffolding the work was built on.

Key points to hit:

- Brief phrase "placing an audience inside a speculative scenario" was the key unlock
- Dunne and Raby: speculative design opens debate, doesn't predict. The three artifacts are prompts not answers
- Kirby: diegetic prototype, artefact from inside the future not describing it from outside. This was the test for every experiment
- *Emily is Away*: interface as a world, being inside rather than watching

- Bandersnatch: active participant versus passive viewer
- Cluely launch video: confidence of format, presents as if the product already exists
- Research was scaffolding, not background. It fed directly into what got built

Section 4: The experiments, in order, with what each one taught me (the bulk of the presentation)

Spoken script / speaker notes:

The first thing I want to say about my experiments is that they were not random. The whole semester followed a double diamond process. Go wide to discover, narrow to define, go wide again to develop, narrow again to deliver. Jed mentioned it early on and I kept coming back to it as a framework every time I felt the pull to commit too early. The experiments were how I moved through that process.

Experiment 1: Learning Claude Code

The first experiment was not capstone-related at all. I built a simple assignment tracker just to understand what Claude Code could actually do before I tried to use it for anything real. What I learned quickly was that being vague does not work. You have to think architecturally before you open the tool. You have to know what you are building, what it should look like, and what you want it to do before you prompt anything. That discipline carried through everything I built after it.

Experiment 2: AI image generation

The second experiment was testing image generation tools to see what they could produce when given concept ideas I had in my head. What I found early was that describing things naturally, the way you would explain an idea to a person, does not translate into these tools at all. The outputs were way off. The better process I found was to first explain the idea to Claude, get it to write a proper prompt, and then take that into the image generator. That extra step made a significant difference in output quality and became a workflow I used for the rest of the semester.

Experiment 3: Storyboarding workflow

Building on the image generation work, I developed a proper storyboarding workflow. Describe what I am after in natural language, pass it to Claude to generate a more refined and detailed prompt, then send that through to an image generation tool. Using Claude as a prompt intermediary rather than prompting directly produced noticeably better results every time. It is a small shift in process but a significant shift in results. This is something I have already started applying outside this course.

Experiment 4: AI video generation

From image generation I moved into video. I tested Higgsfield, Kling, and Gemini to see what they could produce for the Compass film concept. The same lesson from image generation applied here: vague prompts produce confident-looking noise. Specific prompts with clear physical setup, persistent UI, and a raw unaltered environment produced something closer to what I was after. The outputs were not great overall, which I will come back to in the failures section, but the process of testing them gave me a much clearer picture of what the film actually needed to look like.

Experiment 5: Suno music generation

One experiment I ran that did not make it into the blogs explicitly was music. I spent a few hours with Suno trying to do two things: continue on from existing pieces of music, and generate motion graphic style sounds in the vein of what you might hear under a cinematic product launch film. Neither worked the way I needed. Suno does not let you use existing pieces of work as a genuine reference, so continuing someone else's unfinished track was off the table. And the prompts I was giving it for the motion graphic sound direction were not landing. I tried different styles and different approaches for a few hours and could not get anything usable. It is an area I want to come back to, but for now it was a dead end.

Experiment 6: The Compass piano scenario and Python/OpenCV compositing

Coming back to the wearable concept, I needed to prototype what the UI would actually show. I started with a pianist learning a piece because it felt like the most honest version of the question the whole project was asking. What does genuine creative assistance look like at its most basic level? Working out the UI took longer than expected because there is a real difference between an overlay that makes sense on a wearable and one that is just a piano app strapped to someone's face. What I landed on was glowing keys appearing just before you need to press them, note names floating above, a subtle pulse for timing, and a progress bar along the edge. I then built a Python and OpenCV compositing script to process real piano footage with these elements overlaid. Some of the UI elements dropped out in the output. It was partially successful but not clean enough to use in the final film.

Experiment 7: Wearable product landscape research

Before committing to a production approach for the film I spent time mapping the existing wearable product landscape. The references I kept coming back to were Meta Ray-Ban smart glasses for how a consumer wearable can disappear into something people actually want to wear, Brilliant Labs Frame for how minimal and design-forward an AI-integrated device can be, Xreal for the more immersive end, and Snap Spectacles for how a creative-first product gets positioned visually. None of these are exactly what Compass would look like, but mapping where they sit helped me form a clearer picture of the design direction. Something closer to Frame than Vision Pro. Lighter, less imposing, something that signals collaboration rather than dominance.

Experiment 8: Compositing attempts and live footage shoots

I tried two approaches to making the Compass video. The first was letting AI generate the whole thing. It gave me footage but no real control over what it was saying visually. The second was filming it myself and compositing AI-generated UI overlays on top. I also shot

experiments with someone wearing a headset while painting and while taking photos, testing whether the device could feel real on camera. The footage itself was useful and had the right energy. The compositing was still the wall. The UI elements did not sit convincingly and it broke the very illusion the film depended on. I also started orienting myself in After Effects, watching tutorials and understanding what compositing could actually do, but did not get to execution yet.

Experiment 9: The Aotearoa Creative journalism publication

Going wide again in Week 7, I landed on the idea of a fake journalism publication from future Aotearoa. A website that reports on the world as if AI displacement of creative work has already played out. Not explaining the future, just reporting from inside it. I built a basic version in Claude Code with an article, a radio segment, and a masthead. The format worked conceptually. The familiar structure of journalism drops you into the speculative world before you have had time to question it. But it did not survive the Week 9 crit for honest reasons I will come to shortly.

Experiment 10: Muse V1

The Emily is Away game concept became Muse. A choose-your-own-path game built in Claude Code, set inside a fake chat messenger interface in future Aotearoa. The player is a creative working on a project and at every decision point they choose whether to lean on the AI or trust themselves. V1 was a single HTML file, one session, two endings. It was too short and mechanically broken. The choice mechanic did not work properly and the experience did not have enough depth for someone to feel the weight of what it was going for. But it proved the concept was worth pushing.

Experiment 11: Muse V2

After the Week 9 crit the feedback on Muse was clear: not enough depth. I started a completely new Claude Code instance rather than iterating on top of the broken file. The architecture changed from one file to four: index.html, style.css, story.js, and game.js. Separating story.js meant I could edit all the dialogue and branching without touching any code, which matters when the thing I am most likely to keep refining is the writing itself. I expanded from one session to three days with distinct scenes. Day 1 is concept development. Day 2 is on set with a shot that is not working. Day 3 is the edit, the final call on the defining shot of the film. I expanded from two endings to five, each mapped to a leanScore running from 0 to 5. The Purist, The Collaborator, The Negotiator, The Passenger, and The Ghost. The most significant addition was UI corruption. As the leanScore climbs the visual environment shifts with it. At a low score the interface is clean and warm. At a high score the background goes cold blue-black, scanlines appear, a vignette creeps in from the edges. The interface itself reflects how much creative work has shifted to MUSE. The player feels it without being told.

Key points to hit:

- Double diamond was the process framework throughout, diverge before converging

- Claude Code: think architecturally before prompting, vague does not work
- Image generation: natural language fails, Claude as prompt intermediary was the unlock
- Storyboarding workflow: natural language to Claude to refined prompt to generation tool, already using this outside the course
- AI video generation: same lesson as image, specific prompts with clear setup work better
- Suno: music generation failed, both continuing existing tracks and motion graphic sounds, want to come back to it
- Piano scenario and OpenCV: device prompts not performs, that distinction defined the concept, compositing partially worked but not clean enough
- Wearable research: Frame not Vision Pro, lighter, signals collaboration not dominance
- Compositing and live shoots: footage had the right energy, compositing was the wall, After Effects orientation started
- Journalism publication: built it, worked conceptually, did not survive crit
- Muse V1: proved concept, mechanically broken, too shallow
- Muse V2: four file architecture, three days, five endings, leanScore, UI corruption, meaningful before and after
- Every experiment fed the next decision, nothing was random

Section 5: The failures, honest and specific (Suno, compositing, V1 Muse, AI video attempts)

Section 5: The failures, honest and specific

Spoken script / speaker notes:

I want to be upfront about the failures because I think they are actually where the most useful learning happened this semester. Not the things that went wrong in a vague sense, but the specific moments where something broke and I had to figure out why.

Failure 1: AI video generation outputs

The AI video generation attempts for the Compass film did not turn out well. I went into it thinking that if I had a clear enough concept in my head the tools would get me close. They did not. The outputs were either too generic, too uncontrolled, or too far from the specific visual register I was going for. What this taught me is that AI video generation at this stage rewards people who already have a highly specific visual language worked out before they open the tool. If you do not know exactly what you want down to the lighting, the framing, the mood, and the pacing, the tool just produces confident-looking noise. The lesson was not that AI video is bad. It was that my brief to the tool was not specific enough, and that is a skill I need to develop.

Failure 2: The compositing problem

This is the failure I am still sitting with. Every attempt to composite a convincing UI overlay onto real footage failed. The elements did not sit right. They broke the illusion instead of completing it. A fake launch film lives or dies on whether the audience believes what they are seeing, and right now I do not have a clean method for making that work. I tried Python and OpenCV, I tried AI-generated overlays, I started learning After Effects. None of it got to a point I was happy with. What this taught me is that compositing is a genuine craft skill that I underestimated. I came into this thinking it was a technical problem I could solve quickly. It is not. It requires dedicated learning time and proper experimentation in isolation, not as a side task while also building everything else.

Failure 3: Suno and music generation

I spent a few hours with Suno trying to generate music for the Compass film. Two things I wanted to do: continue on from existing tracks and generate motion graphic style sounds. Neither worked. Suno does not allow you to use existing pieces of work as a real reference point, which killed the first idea immediately. And the prompts I was giving it for the second direction were not landing no matter how much I varied them. What this taught me is that music generation requires the same thing image and video generation requires: a highly specific brief. I was going in with a feeling rather than a concrete description, and the tool had nothing real to work with. I also think I just did not spend enough time with it. A few hours is not enough to learn a new tool properly. I should have either committed to it fully or not started at all.

Failure 4: Muse V1

The first version of Muse was too shallow and mechanically broken. The choice mechanic did not function properly, the opening required the player to type a specific phrase to start which felt completely wrong, and there was not enough narrative depth for anyone to feel the weight of what the experience was going for. The concept was right. The execution was not. What this taught me is that interactive experiences require you to map the full architecture before you write a single line of dialogue or code. I jumped into building before I had a clear enough picture of the whole structure, and the output reflected that.

Failure 5: The journalism publication

The journalism publication was a real experiment that I built properly and then let go of at the Week 9 crit. The format worked conceptually. The website looked credible. The articles landed. But when the group asked me directly how passionate I was about journalism as a practice, I had nothing. I could not answer where it goes, what skill I was building, or why I cared about it enough to do it well in 304. Letting it go was the right decision, but it was also a lesson about knowing the difference between a format that is interesting to you intellectually and one you actually have skin in the game with. I was drawn to it because I knew how to execute it, not because it was genuinely mine.

Key points to hit:

- AI video generation: the tool rewards specificity, my brief was not specific enough, that is the real failure
- Compositing: underestimated it as a craft skill, treated it as a quick technical fix, it is not
- Suno: music generation needs a concrete brief not a feeling, also did not spend enough time with it
- Muse V1: jumped into building before mapping the full architecture, output reflected that
- Journalism publication: interesting format but no real investment in it as a practice, knowing the difference between intellectual interest and genuine skin in the game
- Failures were not vague, each one pointed to something specific I need to do differently

Section 6: Skills developed, before and after

One thing this semester clarified for me is how much I actually built in terms of capability. Not just knowledge of tools, but genuine new ways of working that did not exist in my toolkit at the start of semester. I want to be specific about this because I think the before and after is more honest than just saying I learned things.

Before this semester I could: Make commercial films. Work with editors. Storyboard manually. Prompt basic AI chatbots. That was the starting point.

What I can do now that I could not before:

The first is Claude Code. I came in not knowing what it was. I can now set up and run Claude Code to build a functioning interactive web experience from scratch. Muse V2 is evidence of that. A four file architecture, a branching narrative system, a leanScore mechanic, UI corruption tied to player choices. None of that existed in my skillset in Week 1.

The second is AI image generation used properly. Not as a shortcut but as a genuine part of a storyboarding and concept development workflow. The shift was learning to use Claude as a prompt intermediary rather than describing things directly to the generation tool. That single workflow change produced significantly better outputs and is something I have already started using on work outside this course.

The third is AI video generation. I now understand what these tools can and cannot do, what kinds of prompts work, and what the current ceiling is. That understanding is genuinely useful even when the outputs are not good, because knowing the limits of a tool is part of knowing how to use it.

The fourth is the Python and OpenCV compositing script. I built a script from scratch to process real piano footage with UI overlays. It was partially successful. But being able to write and run that script at all is a capability I did not have before.

The fifth is After Effects orientation. I installed it, started learning it, and began understanding what is possible with compositing. I have not executed anything finished in it yet, which I will be honest about. But the orientation work means I am not starting from zero when I go into 304.

The sixth is the prompt intermediary workflow specifically. Describe what I want in natural language, pass it to Claude to generate a refined and detailed prompt, then send that to the generation tool. This applies across image, video, and storyboarding. It is a small process shift and a significant results shift. It is the single most transferable thing I learned this semester.

The seventh is speculative world building. I came in knowing how to make commercial films. I am leaving knowing how to build a world, name it, set it in a specific future Aotearoa context, and design three different formats for putting an audience inside that world rather than explaining it at them. That is a different kind of creative thinking than I had before.

Key points to hit:

- Before: commercial filmmaker, manual storyboarding, basic AI prompting
- Claude Code: can now build functioning interactive web experiences from scratch, Muse V2 is the evidence
- AI image generation: Claude as prompt intermediary was the unlock, already using this outside the course
- AI video generation: understand the limits now, knowing a tool's ceiling is part of knowing how to use it
- Python and OpenCV: built a compositing script from scratch, partially successful
- After Effects: orientation done, not finished execution yet, honest about that
- Prompt intermediary workflow: most transferable single thing learned this semester
- Speculative world building: different kind of creative thinking, came in as commercial filmmaker leaving as someone who can build and inhabit a speculative world

Section 7: Strongest capstone direction plus two genuine alternatives for 304 Week 1

Coming out of this semester I have a clear sense of where I am pointing for 304. I want to be specific about that rather than leaving it open ended.

My strongest direction: The Compass launch film

The speculative launch film for Compass is my primary capstone direction. Not because it is the most complete prototype right now, Muse actually is, but because it sits at the intersection of everything I am most capable of and most genuinely excited to push further. I make films. Pushing that into speculative, world-building territory is not a stretch away from my practice. It is a deepening of it.

What the film needs to be is what the group described at the Week 9 crit. Not a tech demo. Not a product explainer. Something cinematic, minimal, and emotional. Something closer to an Apple Watch campaign than a Kickstarter video. The device is almost beside the point. What the film needs to communicate is the feeling of having something that works with you rather than for you or instead of you.

What I need to refine specifically before this becomes a finished capstone piece: I need to actually design the physical product first. The crit pushed me on this directly. Before committing to a production approach you have to know what the object actually looks like, how it sits on someone's face, what it signals when someone is wearing it. I also need to properly learn After Effects and run dedicated compositing experiments in isolation rather than as a side task. And I need to film proper POV demos to test what the device looks like on camera at different distances and in different creative contexts, piano, painting, photography, and others.

Alternative direction 1: Muse

Muse is my second direction and it is genuinely distinct from the film. It is not a variation on the same idea. Where the film is something you watch and feel, Muse is something you play and inhabit. The interactivity is the point. Every choice the player makes shifts the visual environment and moves the ending. The experience of authorship transferring to the AI is felt rather than explained.

What I would refine in Muse for 304: the writing needs more depth. The branching structure is solid but the dialogue needs to do more emotional work at each decision point. I would also want to test it with people who are not already familiar with the concept, because the real question is whether someone coming to it cold feels the tension the way it is designed to land. And I want to explore whether the interface design can go further, whether the UI corruption can become more visceral and disorienting as the leanScore climbs.

Alternative direction 2: The Aotearoa Creative, revisited

I let the journalism publication go at the Week 9 crit for honest reasons. But I want to name it as a genuine alternative rather than a closed door, because the format itself was not the problem. The problem was my investment in journalism as a practice. If I came back to it in 304 I would reframe it entirely. Not a journalism publication. A broader speculative media world. Multiple formats housed in one place: a publication, a radio segment, a fake job listing for a creative role that no longer exists, a government policy document from 2030. The world building is richer when it comes from multiple angles rather than one format. And if I came back to it with a genuine reason to care about it, which I think I could build through the research, it could be the most conceptually complete of the three.

Key points to hit:

- Strongest direction: Compass launch film, not because most complete but because deepest alignment with practice and passion

- Film needs to feel like Apple Watch campaign not a Kickstarter video, cinematic minimal emotional
- Three specific things to refine before 304: design the physical product first, learn After Effects properly, film POV demos in isolation
- Alternative 1: Muse, genuinely distinct because interactivity is the point, authorship transfer felt not explained
- Muse refinements: deeper dialogue writing, test with people who don't already know the concept, push UI corruption further
- Alternative 2: Aotearoa Creative revisited, not journalism specifically but a broader speculative media world with multiple formats
- All three are distinct directions not variations on the same thing
- Being clear about which one I am pointing at and why

Section 8: What I need to test, build or clarify before 304 starts

This is probably the most practical section of this presentation. Not where I want to go in theory, but what I actually need to do before 304 starts to give myself the best possible entry point into Capstone.

Thing 1: Design the physical product

The biggest gap sitting at the end of this semester is that I never properly designed Compass as a physical object. I thought about it, I researched the wearable landscape, I knew it should sit closer to Frame than Vision Pro. But I never actually designed it. Before I can film anything convincingly I need to know exactly what this object looks like. What it looks like on someone's face. How it sits. What it signals when someone is wearing it in a creative context. That design work needs to happen before any camera goes up.

In hindsight I should have run experiments on actually making or mocking up the physical product earlier in the semester. I held back because I did not want to go too deep into one direction before I had tested the others. That was the right call process-wise but it means I am going into 304 with the physical design still unresolved.

Thing 2: Actually learn After Effects

I installed After Effects and watched tutorials. That is orientation, not skill. Before 304 I need to run dedicated compositing experiments in After Effects in isolation, not as part of trying to also shoot and direct a film at the same time. The compositing problem needs its own focused time. A few sessions where the only goal is understanding what is possible with UI overlays on real footage and getting to a point where I can execute it cleanly.

Thing 3: Film POV demo experiments

Before I commit to a final production approach for the Compass film I need to run proper POV camera tests. What does the device look like on camera at close range, mid range, at a distance. What does the creative context look like from the wearer's perspective. Piano,

painting, photography. These are low stakes experiments that will tell me something specific about how to actually shoot the film. I cannot figure that out by thinking about it. I have to film it.

Thing 4: Test Muse with cold audiences

Muse V2 has only been seen by people who already know what it is and what it is trying to do. Before 304 I want to put it in front of people who have no context for it at all and watch how they respond. Do they feel the tension? Does the UI corruption register? Does the ending land? Those are questions I cannot answer from inside the project. I need external eyes on it.

Thing 5: Revisit music generation properly

Suno was a failure this semester because I did not spend enough time with it and my briefs were not specific enough. Before 304 I want to give it a proper dedicated session with a much more specific brief. Music is going to matter for the Compass film and I would rather figure out whether AI generation can get me there before I start the film than discover mid-production that I need to find another solution.

Thing 6: Clarify the Aotearoa specificity

Jed flagged this as an open question and it is still open. Who specifically are the everyday creatives in Aotearoa I am designing for? The HMW question says everyday creators but that is still broad. Before 304 I need to get more specific about that audience. Who are they, what is their actual relationship with AI tools right now, and how does that shape what the work needs to do for them. That might mean some actual research conversations before the semester starts.

Key points to hit:

- Design the physical product first, everything else depends on it, this is the biggest unresolved gap
- In hindsight should have run physical product experiments earlier, held back to avoid overcommitting, right call process-wise but leaves a gap
- After Effects: orientation is not skill, needs dedicated isolated sessions not as a side task
- POV demo experiments: low stakes tests that answer specific production questions, cannot think your way to those answers
- Test Muse with cold audiences, only been seen by people who already know what it is
- Revisit Suno properly with a specific brief and dedicated time, not a few hours of vague prompting
- Clarify Aotearoa audience specificity, everyday creatives is still too broad, Jed flagged this

Section 9: Positionality and how it shaped decisions

Spoken script / speaker notes:

Positionality is one of those words that can feel like an academic obligation if you let it. I want to talk about it concretely, because I think it actually shaped some of the most important decisions I made this semester in ways I had to actively work against.

The optimism gap

The most significant positionality tension I carried through this semester was the gap between my starting point and the brief's audience. I came in from San Francisco, from rooms full of founders and builders who were excited about AI even when they were afraid of it. That is the world I move in. It is not the world most everyday creatives in Aotearoa move in. For a lot of people here the threat of AI displacement is not speculative. It is immediate and it is real. I had to keep reminding myself of that gap because my natural instinct was to frame everything from the optimist's position.

That gap shaped specific decisions. It is why I designed the Muse game so that neither path is framed as the right one. It is why the Compass film is not an advertisement for AI assistance. It is a portrait of what a different future could look like, and the question of whether you want it is left genuinely open. The work had to be built to function for someone whose relationship to AI is much more anxious than mine, even though my starting point was not anxious at all.

The media fluency advantage and its blind spot

Coming from a filmmaking and media background meant the journalism publication clicked for me immediately. I knew intuitively what a credible publication looks and feels like, what makes a headline land, what makes a byline feel authoritative. That fluency was an advantage in execution. But it was also a blind spot. The same familiarity that made the format feel immediately legible to me made me less likely to question whether it was actually the strongest direction or just the one I felt most equipped to execute. That is exactly what happened. I built it confidently and let it go at the crit when I realised the confidence came from execution comfort, not genuine investment.

The Aotearoa specificity problem

My perspective was shaped heavily by what I saw and heard in San Francisco, and the brief is specifically about Aotearoa and New Zealand stakeholders. Those are not the same context. I named this tension in Week 8 and I want to be honest that I did not fully resolve it this semester. The HMW question is still broader than it should be. Everyday creators in Aotearoa is not specific enough. Jed flagged it. I flagged it myself. Before 304 I need to do the research that narrows that down into something I can actually design for with specificity.

What this changes about how I operate

Going into 304 I want to actively seek out perspectives that are not mine before I start making things. Not as a box to tick but as a genuine design input. The work this semester was built largely from my own research, my own references, my own instincts. That produced interesting things. But it also produced a journalism publication I was confident in for the wrong reasons and a Compass concept that still needs grounding in the actual experiences of Aotearoa creatives rather than just the ones I imagine they have. The positionality work for 304 needs to start with conversations, not concepts.

Key points to hit:

- The optimism gap: came in from SF builder culture, that is not the average Aotearoa creative's relationship with AI
- Optimism gap shaped specific decisions: Muse has no right answer, Compass film leaves the question open, built to work for anxious audiences not just optimistic ones
- Media fluency advantage and blind spot: journalism publication clicked fast because of execution comfort, let it go when I realised that was the wrong reason to be confident
- Aotearoa specificity problem: SF framing versus local context, still not fully resolved, Jed flagged it, I flagged it myself
- What changes: seek perspectives that are not mine before making things, positionality work for 304 starts with conversations not concepts

Section 10: How I operate differently as a designer now

I want to close with something honest rather than something polished. Because I think the most useful thing I can say about how this semester changed me is not about the tools I learned or the prototypes I built. It is about the way I now move through a design problem.

I start from not knowing now

Leo presented his practice-led PhD in Week 10 and the thing that stuck with me was his framing: start from not knowing rather than from a fixed identity. The problem teaches. I came into this semester with a strong sense of what I was going to make. A speculative launch film. That was the plan from Week 1. And I did end up making that, roughly. But the version I am building now looks nothing like the version I described in Week 1 because I let the experiments teach me rather than just confirm what I already thought. That is a different way of working than I had before and I want to keep it.

I now reach for new tools as a first instinct rather than a last resort

Coming into this semester I had a set of tools I was comfortable with and I largely stayed inside them. Video, editing, storytelling. What this semester forced me to do was pick up tools I had never touched before and use them to actually make things, not just learn about them. Claude Code, After Effects, Higgsfield, Kling, Suno, Python and OpenCV. Some of

them worked. Some of them did not. But the act of reaching for something unfamiliar and building with it before I fully understood it changed how I think about my own practice.

The specific shift is this: I used to assess a new tool before committing to it. Watch tutorials, understand the ceiling, decide if it was worth learning. Now I build something with it first and let the output tell me whether it is worth going deeper. That is a faster and more honest way to evaluate a tool than any amount of research beforehand. And it means I am expanding my toolkit through making rather than through planning, which is where I should have been all along.

I now explore further and wider with the tools I use

One thing I would do differently if I had this semester again is go deeper into more tools earlier. I touched a lot of things but I did not always give them enough time to surprise me. Suno was a few hours. After Effects was tutorials. Higgsfield and Kling were early experiments that I moved on from quickly when the outputs were not immediately what I wanted.

What I learned is that the interesting things a tool can do are usually not visible in the first session. You have to push past the obvious use case before you find the thing that is actually useful. I was too quick to move on when something did not work immediately, when I should have stayed longer and pushed harder. That is the exploratory instinct I want to bring into 304. Not just picking up new tools but actually living inside them long enough to find their edges.

I use external pressure as a design tool

The crit sessions were where the most important decisions happened. The journalism publication died at a crit. The Compass concept sharpened at a crit. Muse got its clearest brief from crit feedback. I came in thinking of crits as assessment moments. I am leaving thinking of them as the most valuable part of the process, because they are the only place where someone who does not already know what you were trying to do has to respond to what is actually in front of them. That gap between intention and reception is where the real information lives.

I seek perspectives that are not mine before I start making

This connects back to positionality. My instinct is to research, form a view, and then build. What this semester showed me is that the research needs to include people whose starting point is different from mine, before the concepts form, not after. The Aotearoa specificity problem I flagged is partly a symptom of building from my own references and instincts without enough grounding in the actual experiences of the people I am designing for. That changes in 304.

I treat failure as information not judgment

Every failure this semester had a specific lesson in it. The compositing problem told me I underestimated it as a craft skill. Suno told me I need a concrete brief not a feeling. Muse V1 told me to map the architecture before writing a single line. The journalism publication told

me the difference between execution confidence and genuine investment. None of those lessons came from the things that worked. They all came from the things that did not. I want to keep treating failure that way rather than treating it as something to avoid or minimise.

This semester I came in as a filmmaker. I am leaving as someone who can build interactive experiences, work across multiple speculative formats, use AI tools as genuine creative collaborators rather than shortcuts, and sit with not knowing long enough for the problem to actually teach me something. That feels like a meaningful change.

Key points to hit:

- Start from not knowing: Leo's framing, the problem teaches, came in with a plan and let experiments reshape it
- Reach for new tools first: used to assess before committing, now build first and let the output decide
- Explore further and wider: interesting things in a tool are not visible in the first session, was too quick to move on, need to stay longer and find the edges
- External pressure as design tool: crits are where the most important decisions happened, gap between intention and reception is where real information lives
- Seek perspectives that are not mine before making: positionality work starts with conversations not concepts
- Failure as information not judgment: every failure had a specific lesson, compositing, Suno, Muse V1, journalism publication
- Came in as a filmmaker, leaving as someone who can build interactive experiences, work across speculative formats, use AI as genuine creative collaborator, sit with not knowing